



NETAJI SUBHAS UNIVERSITY OF TECHNOLOGY

High Performance Computing [COCSC18]

Lab File

SUBMITTED BY :-

**Shivam Singla
2019UCO1526
COE-1**

Index

S. No.	Topic	Page No.
1.	Run a basic hello world program using pthreads	3-4
2.	Run a program to find the sum of all elements of an array using 2 processors	5-6
3.	Compute the sum of all the elements of an array using p processors	7-8
4.	Write a program to illustrate basic MPI communication routines	9-10
5.	Design a parallel program for summing up an array, matrix multiplication and show logging and tracing MPI activity	11-14
6.	Write a C program with openMP to implement loop work sharing	15-17
7.	Write a C program with openMP to implement sections work sharing	18-19
8.	Write a program to illustrate process synchronization and collective data movements	20-22

Experiment 1

Run a basic hello world program using pthreads

Code:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(NULL, NULL);
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);
    printf("Hello world! from processor, my rank is %d out of %d
processors\n", world_rank, world_size);
    MPI_Finalize();
}
```

Output:

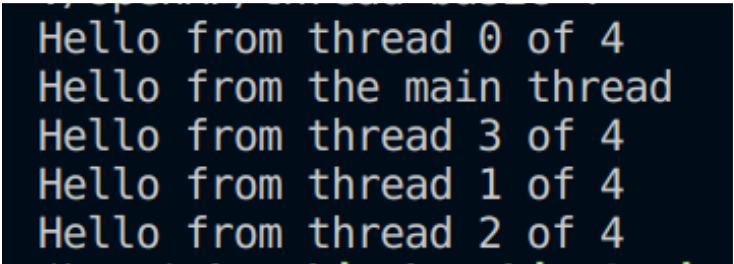
```
C:\Users\sachi\source\repos\Project2\x64\Debug>mpiexec -n 5 ./Project2.exe
Hello world! from processor, my rank is 0 out of 5 processors
Hello world! from processor, my rank is 3 out of 5 processors
Hello world! from processor, my rank is 2 out of 5 processors
Hello world! from processor, my rank is 4 out of 5 processors
Hello world! from processor, my rank is 1 out of 5 processors
```

Using Pthreads

Code:

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
int thread_count;
void* Hello(void* rank);
int main(int argc, char* argv[])
{
    long thread;
    pthread_t* thread_handles;
    thread_count = strtol(argv[1], NULL, 10);
    thread_handles = malloc(thread_count * sizeof(pthread_t));
    for (thread = 0; thread < thread_count; thread++)
        pthread_create(&thread_handles[thread], NULL, Hello, (void*)thread);
    printf("Hello from the main thread\n");
    for (thread = 0; thread < thread_count; thread++)
        pthread_join(thread_handles[thread], NULL);
    free(thread_handles);
    return 0;
}
void* Hello(void* rank)
{
    long my_rank = (long)rank;
    printf("Hello from thread %ld of %d\n", my_rank, thread_count);
    return NULL;
}
```

Output:



```
Hello from thread 0 of 4
Hello from the main thread
Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
```

Experiment 2

Run a program to find the sum of all elements of an array using 2 processors

Code:

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define n 9

int a[] = { 1, 5, 2, 6, 2, 0, 1, 9 };
int a2[1000];

int main(int argc, char* argv[])
{
    int pid, np,
        elements_per_process,
        n_elements_recieved;
    MPI_Status status;
    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    if (pid == 0) {
        int index, i;
        elements_per_process = n / np;
        if (np > 1) {
            for (i = 1; i < np - 1; i++) {
                index = i * elements_per_process;

                MPI_Send(&elements_per_process,
                        1, MPI_INT, i, 0,
                        MPI_COMM_WORLD);
                MPI_Send(&a[index],
                        elements_per_process,
                        MPI_INT, i, 0,
                        MPI_COMM_WORLD);
            }

            index = i * elements_per_process;
            int elements_left = n - index;

            MPI_Send(&elements_left,
                    1, MPI_INT,
                    i, 0,
                    MPI_COMM_WORLD);
            MPI_Send(&a[index],
                    elements_left,
```

```

        MPI_INT, i, 0,
        MPI_COMM_WORLD);
    }

    int sum = 0;
    for (i = 0; i < elements_per_process; i++)
        sum += a[i];

    int tmp;
    for (i = 1; i < np; i++) {
        MPI_Recv(&tmp, 1, MPI_INT,
                MPI_ANY_SOURCE, 0,
                MPI_COMM_WORLD,
                &status);
        int sender = status.MPI_SOURCE;

        sum += tmp;
    }

    printf("Sum of array is : %d\n", sum);
}
else {
    MPI_Recv(&n_elements_recieved,
            1, MPI_INT, 0, 0,
            MPI_COMM_WORLD,
            &status);

    MPI_Recv(&a2, n_elements_recieved,
            MPI_INT, 0, 0,
            MPI_COMM_WORLD,
            &status);

    int partial_sum = 0;
    for (int i = 0; i < n_elements_recieved; i++)
        partial_sum += a2[i];

    MPI_Send(&partial_sum, 1, MPI_INT,
            0, 0, MPI_COMM_WORLD);
}

MPI_Finalize();

return 0;
}

```

Output:

```

C:\Users\sachi\source\repos\Project2\x64\Debug>mpiexec ./Project2.exe
Sum of array provided is : 26

```

Experiment 3

Compute the sum of all the elements of an array using p processors

Code:

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#define n 8
int a[] = { 1, 5, 2, 6, 2, 0, 1, 9 };
// Temporary array for other processes
int b[1000];
int main(int argc, char* argv[])
{
    int process_id, no_of_process,
        elements_per_process,
        n_elements_recieved;
    MPI_Status status;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &process_id);
    MPI_Comm_size(MPI_COMM_WORLD, &no_of_process);
    // For process 0
    if (process_id == 0) {
        int index, i;
        elements_per_process = n / no_of_process;
        if (no_of_process > 1) {
            for (i = 1; i < no_of_process - 1; i++) {
                index = i * elements_per_process;
                MPI_Send(&elements_per_process, 1, MPI_INT, i, 0,
MPI_COMM_WORLD);
                MPI_Send(&a[index], elements_per_process, MPI_INT, i, 0,
MPI_COMM_WORLD);
            }
            // last process adds remaining elements
            index = i * elements_per_process;
            int elements_left = n - index;
            MPI_Send(&elements_left, 1, MPI_INT, i, 0, MPI_COMM_WORLD);
            MPI_Send(&a[index], elements_left, MPI_INT, i, 0,
MPI_COMM_WORLD);
        }
        // sum by process 0
        int sum = 0;
        for (i = 0; i < elements_per_process; i++)
            sum += a[i];
        printf("Sum by this Slave Process is %d = %d\n", process_id, sum);
        // partial sums from other processes
        int tmp;
        for (i = 1; i < no_of_process; i++) {
            MPI_Recv(&tmp, 1, MPI_INT, MPI_ANY_SOURCE, 0, MPI_COMM_WORLD,
&status);
            int sender = status.MPI_SOURCE;
```

```

        sum += tmp;
    }
    // prints the final sum of array
    printf("Final Sum of array by the Master Process is : %d\n", sum);
}
// Other processes
else {
    MPI_Recv(&n_elements_recieved, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
&status);
    MPI_Recv(&b, n_elements_recieved, MPI_INT, 0, 0, MPI_COMM_WORLD,
&status);
    int partial_sum = 0;
    for (int i = 0; i < n_elements_recieved; i++)
        partial_sum += b[i];
    printf("Sum by process %d = %d\n", process_id, partial_sum);
    MPI_Send(&partial_sum, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
}
MPI_Finalize();
return 0;
}

```

Output :

```

C:\Users\sachi\source\repos\Project2\x64\Debug>mpiexec -n 4 ./Project2.exe
Sum by process 3 = 10
Sum by process 2 = 2
Sum by process 1 = 8
Sum by this Slave Process is 0 = 6
Final Sum of array by the Master Process is : 26

```


Experiment 4

Write a program to illustrate basic MPI communication routines

Code:

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char** argv)
{
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
    // COMM_WORLD is the communicator world
    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    // Get the name of the processor
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int name_len;
    MPI_Get_processor_name(processor_name, &name_len);
    printf("Hello world from process %s, rank %d out of %d processes\n\n",
           processor_name, world_rank, world_size);
    if (world_rank == 0)
    {
        char message[] = "Shivam";
        MPI_Send(message, 6, MPI_CHAR, 1, 0, MPI_COMM_WORLD);
    }
    else
    {
        char message[6];
        MPI_Recv(message, 6, MPI_CHAR, 0, 0, MPI_COMM_WORLD,
MPI_STATUS_IGNORE);
        printf("Message Successfully Received\n");
        printf("Message Recieved : %s\n", message);
    }
    MPI_Finalize();
    return 0;
}
```

Output:

```
C:\Users\sachi\source\repos\Project2\x64\Debug>mpiexec -n 2 ./Project2.exe
Hello world from process LAPTOP-MSQNOU3G, rank 0 out of 2 processes
Hello world from process LAPTOP-MSQNOU3G, rank 1 out of 2 processes
Message received!
Message is : Shivam
```

Experiment 5

Design a parallel program for summing up an array, matrix multiplication and show logging and tracing MPI activity

Code:

```
#include<stdio.h>
#include<iostream>
#include "mpi.h"
#define NUM_ROWS_A 8
#define NUM_COLUMNS_A 10
#define NUM_ROWS_B 10
#define NUM_COLUMNS_B 8
#define MASTER_TO_SLAVE_TAG 1 //tag for messages sent from master to slaves
#define SLAVE_TO_MASTER_TAG 4 //tag for messages sent from slaves to master
void create_matrix();
void printArray();
int rank;
int size;
int i, j, k;
double A[NUM_ROWS_A][NUM_COLUMNS_A];
double B[NUM_ROWS_B][NUM_COLUMNS_B];
double result[NUM_ROWS_A][NUM_COLUMNS_B];
int low_bound; //low bound of the number of rows of [A] allocated to a slave
int upper_bound; //upper bound of the number of rows of [A] allocated to a slave
int portion; //portion of the number of rows of [A] allocated to a slave
MPI_Status status; // store status of a MPI_Recv
MPI_Request request; //capture request of a MPI_Send
int main(int argc, char* argv[])
{
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    if (rank == 0)
    { // master process
        create_matrix();
        for (i = 1; i < size; i++)
        {
            portion = (NUM_ROWS_A / (size - 1)); // portion without master
            low_bound = (i - 1) * portion;
            if (((i + 1) == size) && ((NUM_ROWS_A % (size - 1)) != 0))
            { //if rows of [A] cannot be equally divided among slaves
                upper_bound = NUM_ROWS_A; //last slave gets all the
remaining rows
            }
            else {
                upper_bound = low_bound + portion; //rows of [A]
are equally divisible among slaves
            }
            MPI_Send(&low_bound, 1, MPI_INT, i, MASTER_TO_SLAVE_TAG,
MPI_COMM_WORLD);
            MPI_Send(&upper_bound, 1, MPI_INT, i, MASTER_TO_SLAVE_TAG + 1,
```

```

        MPI_COMM_WORLD);
        MPI_Send(&A[low_bound][0], (upper_bound - low_bound) *
NUM_COLUMNS_A,
        MPI_DOUBLE, i, MASTER_TO_SLAVE_TAG + 2, MPI_COMM_WORLD);
    }
}
//broadcast [B] to all the slaves
MPI_Bcast(&B, NUM_ROWS_B * NUM_COLUMNS_B, MPI_DOUBLE, 0, MPI_COMM_WORLD);
/* Slave process*/
if (rank > 0)
{
    MPI_Recv(&low_bound, 1, MPI_INT, 0, MASTER_TO_SLAVE_TAG,
MPI_COMM_WORLD,
        &status);
    MPI_Recv(&upper_bound, 1, MPI_INT, 0, MASTER_TO_SLAVE_TAG + 1,
        MPI_COMM_WORLD, &status);
    MPI_Recv(&A[low_bound][0], (upper_bound - low_bound) * NUM_COLUMNS_A,
        MPI_DOUBLE, 0, MASTER_TO_SLAVE_TAG + 2, MPI_COMM_WORLD,
&status);
    printf("Process %d calculating for rows %d to %d of Matrix A\n", rank,
        low_bound, upper_bound);
    for (i = low_bound; i < upper_bound; i++)
    {
        for (j = 0; j < NUM_COLUMNS_B; j++)
        {
            for (k = 0; k < NUM_ROWS_B; k++)
            {
                result[i][j] += (A[i][k] * B[k][j]);
            }
        }
    }
    MPI_Send(&low_bound, 1, MPI_INT, 0, SLAVE_TO_MASTER_TAG,
MPI_COMM_WORLD);
    MPI_Send(&upper_bound, 1, MPI_INT, 0, SLAVE_TO_MASTER_TAG + 1,
        MPI_COMM_WORLD);
    MPI_Send(&result[low_bound][0], (upper_bound - low_bound) *
NUM_COLUMNS_B,
        MPI_DOUBLE, 0, SLAVE_TO_MASTER_TAG + 2, MPI_COMM_WORLD);
}
/* master gathers processed work*/
if (rank == 0) {
    for (i = 1; i < size; i++) {
        MPI_Recv(&low_bound, 1, MPI_INT, i, SLAVE_TO_MASTER_TAG,
MPI_COMM_WORLD,
            &status);
        MPI_Recv(&upper_bound, 1, MPI_INT, i, SLAVE_TO_MASTER_TAG + 1,
            MPI_COMM_WORLD, &status);
        MPI_Recv(&result[low_bound][0], (upper_bound - low_bound)
*
            NUM_COLUMNS_B, MPI_DOUBLE, i, SLAVE_TO_MASTER_TAG +
2, MPI_COMM_WORLD, &status);
    }
    printArray();
}
MPI_Finalize();
return 0;
}
void create_matrix()

```

```

{
    for (i = 0; i < NUM_ROWS_A; i++) {
        for (j = 0; j < NUM_COLUMNS_A; j++) {
            A[i][j] = i + j;
        }
    }
    for (i = 0; i < NUM_ROWS_B; i++) {
        for (j = 0; j < NUM_COLUMNS_B; j++) {
            B[i][j] = i * j;
        }
    }
}

void printArray()
{
    printf("Given matrix A is: \n");
    for (i = 0; i < NUM_ROWS_A; i++) {
        printf("\n");
        for (j = 0; j < NUM_COLUMNS_A; j++)
            printf("%8.2f ", A[i][j]);
    }
    printf("\n\n\n");
    printf("Given matrix B is: \n");
    for (i = 0; i < NUM_ROWS_B; i++) {
        printf("\n");
        for (j = 0; j < NUM_COLUMNS_B; j++)
            printf("%8.2f ", B[i][j]);
    }
    printf("\n\n\n");
    printf("Final Multiplied Matrix is: \n");
    for (i = 0; i < NUM_ROWS_A; i++) {
        printf("\n");
        for (j = 0; j < NUM_COLUMNS_B; j++)
            printf("%8.2f ", result[i][j]);
    }
    printf("\n\n\n");
}

```

Output:

```
C:\Users\sachi\source\repos\Project2\x64\Debug>mpiexec -n 4 ./Project2.exe
Given matrix A is:
```

0.00	1.00	2.00	3.00	4.00	5.00	6.00	7.00	8.00	9.00
1.00	2.00	3.00	4.00	5.00	6.00	7.00	8.00	9.00	10.00
2.00	3.00	4.00	5.00	6.00	7.00	8.00	9.00	10.00	11.00
3.00	4.00	5.00	6.00	7.00	8.00	9.00	10.00	11.00	12.00
4.00	5.00	6.00	7.00	8.00	9.00	10.00	11.00	12.00	13.00
5.00	6.00	7.00	8.00	9.00	10.00	11.00	12.00	13.00	14.00
6.00	7.00	8.00	9.00	10.00	11.00	12.00	13.00	14.00	15.00
7.00	8.00	9.00	10.00	11.00	12.00	13.00	14.00	15.00	16.00

```
Given matrix B is:
```

0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	2.00	3.00	4.00	5.00	6.00	7.00
0.00	2.00	4.00	6.00	8.00	10.00	12.00	14.00
0.00	3.00	6.00	9.00	12.00	15.00	18.00	21.00
0.00	4.00	8.00	12.00	16.00	20.00	24.00	28.00
0.00	5.00	10.00	15.00	20.00	25.00	30.00	35.00
0.00	6.00	12.00	18.00	24.00	30.00	36.00	42.00
0.00	7.00	14.00	21.00	28.00	35.00	42.00	49.00
0.00	8.00	16.00	24.00	32.00	40.00	48.00	56.00
0.00	9.00	18.00	27.00	36.00	45.00	54.00	63.00

```
Given matrix B is:
```

0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
0.00	1.00	2.00	3.00	4.00	5.00	6.00	7.00
0.00	2.00	4.00	6.00	8.00	10.00	12.00	14.00
0.00	3.00	6.00	9.00	12.00	15.00	18.00	21.00
0.00	4.00	8.00	12.00	16.00	20.00	24.00	28.00
0.00	5.00	10.00	15.00	20.00	25.00	30.00	35.00
0.00	6.00	12.00	18.00	24.00	30.00	36.00	42.00
0.00	7.00	14.00	21.00	28.00	35.00	42.00	49.00
0.00	8.00	16.00	24.00	32.00	40.00	48.00	56.00
0.00	9.00	18.00	27.00	36.00	45.00	54.00	63.00

```
Final Multiplied Matrix is:
```

0.00	285.00	570.00	855.00	1140.00	1425.00	1710.00	1995.00
0.00	330.00	660.00	990.00	1320.00	1650.00	1980.00	2310.00
0.00	375.00	750.00	1125.00	1500.00	1875.00	2250.00	2625.00
0.00	420.00	840.00	1260.00	1680.00	2100.00	2520.00	2940.00
0.00	465.00	930.00	1395.00	1860.00	2325.00	2790.00	3255.00
0.00	510.00	1020.00	1530.00	2040.00	2550.00	3060.00	3570.00
0.00	555.00	1110.00	1665.00	2220.00	2775.00	3330.00	3885.00
0.00	600.00	1200.00	1800.00	2400.00	3000.00	3600.00	4200.00

```
Process 2 calculating for rows 2 to 4 of Matrix A
```

```
Process 1 calculating for rows 0 to 2 of Matrix A
```

```
Process 3 calculating for rows 4 to 8 of Matrix A
```

Experiment 6

Write a C program with openMP to implement loop work sharing

Code:

```
#include <omp.h>
#include <stdio.h>
#include <iostream>

using namespace std;
void reset_freq(int* freq, int THREADS)
{
    for (int i = 0; i < THREADS; i++)
        freq[i] = 0;
}
int main(int* argc, char** argv)
{
    int n, THREADS, i;
    printf("Enter the number of iterations :");
    scanf_s("%d", &n);
    printf("Enter the number of threads (max 8): ");
    scanf_s("%d", &THREADS);
    int freq[6];
    reset_freq(freq, THREADS);
    // simple parallel for with unequal iterations
#pragma omp parallel for num_threads(THREADS)
    for (i = 0; i < n; i++)
    {
        // printf("Thread num %d executing iter %d\n", omp_get_thread_num(),
i);
        freq[omp_get_thread_num()]++;
    }
#pragma omp barrier
    printf("\nIn default scheduling, we have the following thread distribution :-
\n");
    for (int i = 0; i < THREADS; i++)
    {
        printf("Thread No. %d : %d iters\n", i, freq[i]);
    }
    // using static scheduling
    int CHUNK;
    printf("\nUsing static scheduling...\n");
    printf("Enter the chunk size :");
    scanf_s("%d", &CHUNK);
    // using a static, round robin schedule for the loop iterations
    reset_freq(freq, THREADS);
    // useful when the workload is ~ same across each thread, not when otherwise
#pragma omp parallel for num_threads(THREADS) schedule(static, CHUNK)
    for (i = 0; i < n; i++)
    {
```

```

        // printf("Thread num %d executing iter %d\n", omp_get_thread_num(),
i);
        freq[omp_get_thread_num()]++;
    }
#pragma omp barrier
    printf("\nIn static scheduling, we have the following thread distribution :-
\n");
    for (int i = 0; i < THREADS; i++)
    {
        printf("Thread No. %d : %d iterations\n", i, freq[i]);
    }
    // auto scheduling depending on the compiler
    printf("\nUsing automatic scheduling...\n");
    reset_freq(freq, THREADS);
#pragma omp parallel for num_threads(THREADS) schedule(static)
    for (i = 0; i < n; i++)
    {
        // printf("Thread num %d executing iter %d\n", omp_get_thread_num(),
i);
        freq[omp_get_thread_num()]++;
    }
#pragma omp barrier
    printf("In auto scheduling, we have the following thread distribution :-
\n");
    for (int i = 0; i < THREADS; i++)
    {
        printf("Thread No. %d : %d iters\n", i, freq[i]);
    }
    return 0;
}

```


Output:

```
Enter the number of iterations :4
Enter the number of threads (max 8): 5

In default scheduling, we have the following thread distribution :-
Thread No. 0 : 1 iters
Thread No. 1 : 1 iters
Thread No. 2 : 1 iters
Thread No. 3 : 1 iters
Thread No. 4 : 0 iters

Using static scheduling...
Enter the chunk size :2

In static scheduling, we have the following thread distribution :-
Thread No. 0 : 2 iterations
Thread No. 1 : 2 iterations
Thread No. 2 : 0 iterations
Thread No. 3 : 0 iterations
Thread No. 4 : 0 iterations

Using automatic scheduling...
In auto scheduling, we have the following thread distribution :-
Thread No. 0 : 1 iters
Thread No. 1 : 1 iters
Thread No. 2 : 1 iters
Thread No. 3 : 1 iters
Thread No. 4 : 0 iters
```


Experiment 7

Write a C program with openMP to implement loop work sharing

Code:

```
#include <omp.h>
#include <stdio.h>
int main(int* argc, char** argv)
{ // invocation of the main program
  // use the fopenmp flag for compiling
  int num_threads, THREAD_COUNT = 4;
  int thread_ID;
  int section_sizes[4] = {
    0, 100, 200, 300 };
  printf("Implementing Work load sharing of threads...\n");
#pragma omp parallel private(thread_ID) num_threads(THREAD_COUNT)
  {
    // private means each thread will have a private variable
    // thread_ID
    thread_ID = omp_get_thread_num();
    printf("I am thread number %d!\n", thread_ID);
    int value_count = 0;
    if (thread_ID > 0)
    {
      int work_load = section_sizes[thread_ID];
      // each thread has a different section size
      for (int i = 0; i < work_load; i++)
        value_count++;
      printf("Total Number of values computed are : %d\n",
value_count);
    }
#pragma omp barrier
    if (thread_ID == 0)
    {
      printf("The Total number of threads are : %d",
omp_get_num_threads());
    }
    return 0;
  }
}
```

Output:

 Microsoft Visual Studio Debug Console

```
Implementing Work load sharing of threads...  
I am thread number 0!  
I am thread number 2!  
I am thread number 3!  
Total Number of values computed are : 300  
I am thread number 1!  
Total Number of values computed are : 100  
Total Number of values computed are : 200  
The Total number of threads are : 4
```

Experiment 8

Write a program to illustrate process synchronization and collective data movements

Code:

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
int thread_count; // this global variable is shared by all threads
// compiling information -
// gcc name_of_file.c -o name_of_exe -lpthread (link p thread)
// necessary for referencing in the thread
struct arguments
{
    int size;
    int* arr1;
    int* arr2;
    int* dot;
};
// function to parallelize`
void* add_into_one(void* arguments);
// util
void print_vector(int n, int* arr)
{
    printf("[ ");
    for (int i = 0; i < n; i++)
        printf("%d ", arr[i]);
    printf("] \n");
}
// main driver function of the program
int main(int argc, char* argv[])
{
    long thread;
    // /* Use long in case of a 64-bit system */
    pthread_t* thread_handles;
    thread_count = 2; // using 2 threads only
    // get the thread handles equal to total num
    // of threads
    thread_handles = malloc(thread_count * sizeof(pthread_t));
    printf("Enter the size of the vectors : ");
    int n;
    scanf("%d", &n);
    printf("Enter the max_val of the vectors : ");
    int max_val;
    scanf("%d", &max_val);
```

```

struct arguments* args[2]; // array of pointer to structure
// each element is a pointer
for (int i = 0; i < 2; i++)
{
    // allocate for the struct
    args[i] = malloc(sizeof(struct arguments) * 1);
    // allocate for the arrays
    args[i]->size = n;
    args[i]->arr1 = malloc(sizeof(int) * n);
    args[i]->arr2 = malloc(sizeof(int) * n);
    args[i]->dot = malloc(sizeof(int) * n);
    for (int j = 0; j < n; j++)
    {
        args[i]->arr1[j] = rand() % max_val;
        args[i]->arr2[j] = rand() % max_val;
    }
}
printf("Vectors are : \n");
print_vector(n, args[0]->arr1);
print_vector(n, args[0]->arr2);
print_vector(n, args[1]->arr1);
print_vector(n, args[1]->arr2);
int result[n];
memset(result, 0, n * sizeof(int));
// note : we need to manually startup our threads
// for a particular function which we want to execute in
// the thread
for (thread = 0; thread < thread_count; thread++)
{
    printf("Multiplying %d and %d with thread %d...\n", thread + 1,
thread + 2,
thread);
    pthread_create(&thread_handles[thread], NULL, add_into_one,
(void*)args[thread]);
}
printf("Hello from the main thread\n");
// wait for completion
for (thread = 0; thread < thread_count; thread++)
    pthread_join(thread_handles[thread], NULL);
for (int i = 0; i < 2; i++)
{
    printf("Multiplication for vector %d and %d \n", i + 1, i + 2);
    print_vector(n, args[i]->dot);
    printf("\n");
}
free(thread_handles);
// now compute the summation of results
for (int i = 0; i < n; i++)
{
    result[i] = args[0]->dot[i] + args[1]->dot[i];
    printf("Result is : \n");
    print_vector(n, result);
}
return 0;
}
void* add_into_one(void* argument)
{
    // de reference the argument

```

```

    struct arguments* args = argument;
    // compute the dot product into the
    // array dot
    int n = args->size;
    for (int i = 0; i < n; i++)
        args->dot[i] = args->arr1[i] * args->arr2[i];
    return NULL;
}

```

Output:

```

Enter the size of the vectors : 6
Enter the max_val of the vectors : 4
Vectors are :
[ 3 1 1 2 1 2 ]
[ 2 3 3 0 1 3 ]
[ 2 3 0 0 3 3 ]
[ 3 2 2 0 0 1 ]
Multiplying 1 and 2 with thread 0...
Multiplying 2 and 3 with thread 1...
Hello from the main thread
Multiplication for vector 1 and 2
[ 6 3 3 0 1 6 ]

Multiplication for vector 2 and 3
[ 6 6 0 0 0 3 ]

Result is :
[ 12 9 3 0 1 9 ]

```